

# Time Series Assignment 1

Arindam Patra

03/11/2021

Name: Arindam Patra, Roll: MDS202008

## Importing libraries

```
knitr::opts_chunk$set(echo = TRUE, warning = FALSE, tidy = TRUE)
library(tseries)
```

```
## Warning: package 'tseries' was built under R version 4.0.5
```

```
## Registered S3 method overwritten by 'quantmod':
```

```
##   method      from
```

```
##   as.zoo.data.frame zoo
```

```
library(forecast)
```

```
## Warning: package 'forecast' was built under R version 4.0.5
```

```
library(moments)
```

## Importing Data

Monthly Electric Production data from 1985-2017


```
df <- read.csv("time_series/Electric_Production.csv", row.names="DATE")
```

```
data = ts(df)
```

```
Time <- as.Date(rownames(df), '%m/%d/%Y')
```

```
plot(Time, data, xlab = 'Date', ylab = 'Electric Production',
      main = 'Electric Production vs Time', type = 'l')
```

## Electric Production vs Time



## Descriptive Statistics

```
m = mean(data)
s = sd(data)
print(paste("Minimum = ", min(data)))

## [1] "Minimum = 55.3151"
print(paste("Maximum = ", max(data)))

## [1] "Maximum = 129.4048"
print(paste("Range = ", max(data) - min(data)))

## [1] "Range = 74.0897"
print(paste("Mean = ", m))

## [1] "Mean = 88.8472176322418"
print(paste("Standard deviation = ", s))

## [1] "Standard deviation = 15.3878336647309"
print(paste("1st Quartile = ", quantile(data, 0.25)))

## [1] "1st Quartile = 77.1052"
```

```

print(paste("2nd Quartile (Median) = ", quantile(data, 0.5)))

## [1] "2nd Quartile (Median) = 89.7795"
print(paste("3rd Quartile = ", quantile(data, 0.75)))

## [1] "3rd Quartile = 100.5244"
print(paste("Inter quartile range = ", IQR(data)))

## [1] "Inter quartile range = 23.4192"
print(paste("Skewness = ", skewness(data)))

## [1] "Skewness = -0.0728191541455781"
print(paste("Kurtosis = ", kurtosis(data)))

## [1] "Kurtosis = 2.2994360495705"
x = round(((length(data[data > m - (3 * s) & data < m + (3 * s)]) * 100)/length(data)),
3)
print(paste("Percentage of observations in 3-sigma deviation of the mean = ", x, "%"))

## [1] "Percentage of observations in 3-sigma deviation of the mean = 100 %"

```

Since all the data lies between 3 SD, so the data doesn't have any outlier.

## Stationarity Test

Let's test the stationarity of this time series. We will test it using 3 tests - ADF, KPSS and PP Test.

```

adf = adf.test(data)
adf

##
## Augmented Dickey-Fuller Test
##
## data: data
## Dickey-Fuller = -5.139, Lag order = 7, p-value = 0.01
## alternative hypothesis: stationary

kpss = kpss.test(data)
kpss

##
## KPSS Test for Level Stationarity
##
## data: data
## KPSS Level = 6.3058, Truncation lag parameter = 5, p-value = 0.01

pp = PP.test(data)
pp

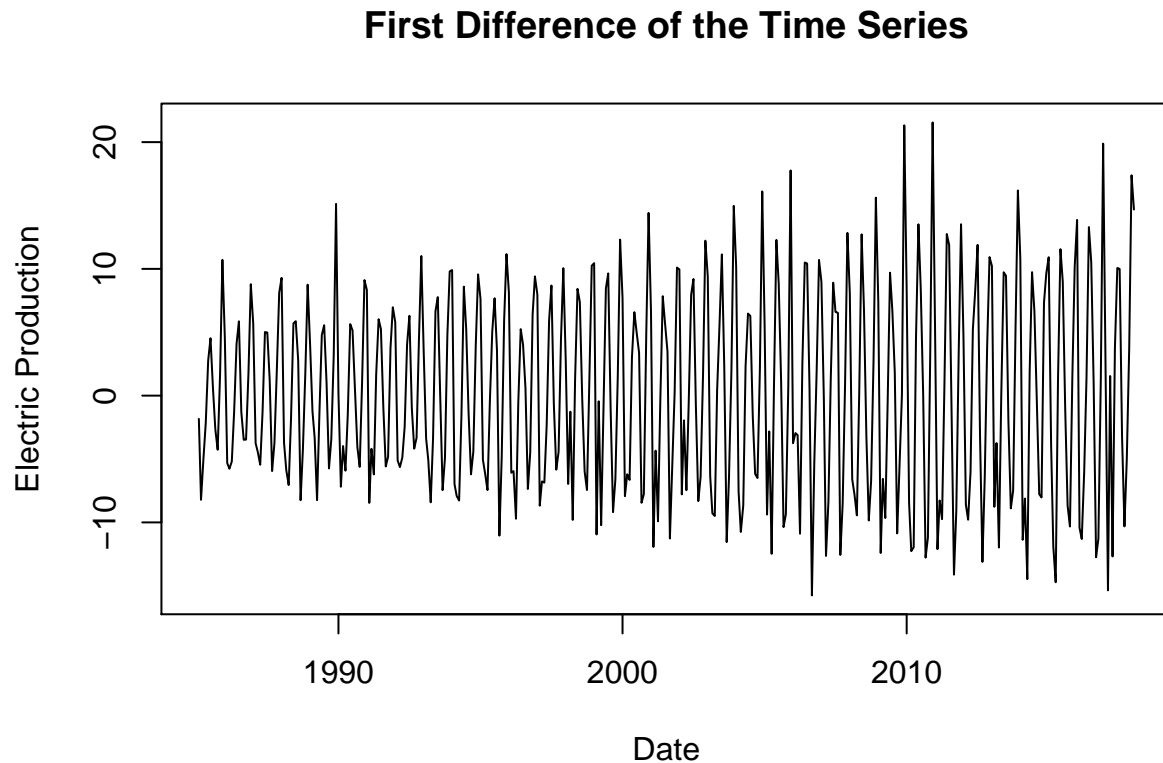
##
## Phillips-Perron Unit Root Test
##
## data: data
## Dickey-Fuller = -8.2573, Truncation lag parameter = 5, p-value = 0.01

```

Here ADF and PP test are agreed for stationarity of this data but KPSS is saying non-stationary. So let's do the first difference of the data.

## First Difference

```
first_diff <- diff(data)
plot(Time[-1],first_diff, xlab = 'Date',ylab = 'Electric Production',
     main = 'First Difference of the Time Series',type = 'l')
```



```
adf = adf.test(first_diff)
adf
```

```
##
## Augmented Dickey-Fuller Test
##
## data: first_diff
## Dickey-Fuller = -9.6447, Lag order = 7, p-value = 0.01
## alternative hypothesis: stationary
```

```
kpss = kpss.test(first_diff)
kpss
```

```
##
## KPSS Test for Level Stationarity
##
## data: first_diff
## KPSS Level = 0.078012, Truncation lag parameter = 5, p-value = 0.1
```

```
pp = PP.test(first_diff)
pp
```

```
##
##  Phillips-Perron Unit Root Test
##
## data:  first_diff
## Dickey-Fuller = -15.573, Truncation lag parameter = 5, p-value = 0.01
```

Ok! Here all of them agrees with the stationarity of this time series data. Therefore we can fit this data to our model.

## Train test split

Keeping last 20 observations as test data and rest of them as train data.

```
test_size <- 20
n <- length(first_diff)
train_size <- n - test_size

train_data <- first_diff[c(1 : train_size)]
test_data <- first_diff[c((train_size+1) : n)]
```

```
### Total data points
length(first_diff)
```

```
## [1] 396
```

```
### Training Data Size
length(train_data)
```

```
## [1] 376
```

```
### Test Data Size
length(test_data)
```

```
## [1] 20
```

## Modelling

At first we will build two optimized model. 1 - Using AIC, 2 - Using BIC

Using AIC

```
aic_model <- auto.arima(train_data, trace= TRUE, d= 0, max.p = 10, max.q = 10,
                        ic ="aic", approximation = FALSE)
```

```
##
## ARIMA(2,0,2) with non-zero mean : 2082.145
## ARIMA(0,0,0) with non-zero mean : 2593.348
## ARIMA(1,0,0) with non-zero mean : 2536.677
## ARIMA(0,0,1) with non-zero mean : 2470.228
## ARIMA(0,0,0) with zero mean : 2591.363
## ARIMA(1,0,2) with non-zero mean : 2345.833
## ARIMA(2,0,1) with non-zero mean : 2089.945
## ARIMA(3,0,2) with non-zero mean : 2076.4
```

```
## ARIMA(3,0,1) with non-zero mean : 2076.408
## ARIMA(4,0,2) with non-zero mean : 2064.957
## ARIMA(4,0,1) with non-zero mean : 2071.16
## ARIMA(5,0,2) with non-zero mean : 2049.122
## ARIMA(5,0,1) with non-zero mean : 2072.437
## ARIMA(6,0,2) with non-zero mean : Inf
## ARIMA(5,0,3) with non-zero mean : Inf
## ARIMA(4,0,3) with non-zero mean : Inf
## ARIMA(6,0,1) with non-zero mean : 2072.464
## ARIMA(6,0,3) with non-zero mean : Inf
## ARIMA(5,0,2) with zero mean : 2060.629
##
## Best model: ARIMA(5,0,2) with non-zero mean
```

```
aic_model
```

```
## Series: train_data
## ARIMA(5,0,2) with non-zero mean
##
## Coefficients:
##          ar1      ar2      ar3      ar4      ar5      ma1      ma2      mean
##          0.6928 -0.3857 -0.2180 -0.0489  0.2877 -0.8282 -0.0991  0.1095
## s.e.    0.1446   0.1248   0.0998   0.0734  0.0573   0.1451   0.1311  0.0208
##
## sigma^2 estimated as 13.12: log likelihood=-1015.56
## AIC=2049.12 AICc=2049.61 BIC=2084.49
```

For AIC, Best model is ARMA(5,2).

Using BIC

```
bic_model <- auto.arima(train_data, trace= TRUE, d=0, max.p = 10, max.q = 10,
                        ic ="bic", approximation = FALSE)
```

```
##
## ARIMA(2,0,2) with non-zero mean : 2105.723
## ARIMA(0,0,0) with non-zero mean : 2601.207
## ARIMA(1,0,0) with non-zero mean : 2548.466
## ARIMA(0,0,1) with non-zero mean : 2482.016
## ARIMA(0,0,0) with zero mean : 2595.293
## ARIMA(1,0,2) with non-zero mean : 2365.481
## ARIMA(2,0,1) with non-zero mean : 2109.593
## ARIMA(3,0,2) with non-zero mean : 2103.907
## ARIMA(3,0,1) with non-zero mean : 2099.985
## ARIMA(3,0,0) with non-zero mean : 2166.644
## ARIMA(4,0,1) with non-zero mean : 2098.667
## ARIMA(4,0,0) with non-zero mean : 2093.737
## ARIMA(5,0,0) with non-zero mean : 2098.297
## ARIMA(5,0,1) with non-zero mean : 2103.874
## ARIMA(4,0,0) with zero mean : 2089.178
## ARIMA(3,0,0) with zero mean : 2161.152
## ARIMA(5,0,0) with zero mean : 2093.927
## ARIMA(4,0,1) with zero mean : 2094.259
## ARIMA(3,0,1) with zero mean : 2102.707
## ARIMA(5,0,1) with zero mean : 2099.486
```

```
##
## Best model: ARIMA(4,0,0) with zero mean
bic_model

## Series: train_data
## ARIMA(4,0,0) with zero mean
##
## Coefficients:
##          ar1      ar2      ar3      ar4
##      -0.0231  -0.4336  -0.4910  -0.4355
## s.e.   0.0466   0.0390   0.0388   0.0467
##
## sigma^2 estimated as 14.03:  log likelihood=-1029.76
## AIC=2069.53   AICc=2069.69   BIC=2089.18
```

For BIC, Best model is ARMA(4,0).

## Choosing BEST Model

We will check MSE for both of these models. For whichever model the MSE is lowest, we will choose that.

```
pred_aic<- forecast(aic_model, h = 20)
aic_forecast<- pred_aic$mean
mse_aic = sum((test_data - aic_forecast)^2)/length(test_data)
print(paste("MSE for AIC Model = ", mse_aic))
```

```
## [1] "MSE for AIC Model = 37.4904858003665"
```

```
pred_bic <- forecast(bic_model, h = 20)
bic_forecast<- pred_bic$mean
mse_bic = sum((test_data - bic_forecast)^2)/length(test_data)
print(paste("MSE for BIC Model = ", mse_bic))
```


```
## [1] "MSE for BIC Model = 38.2007743700181"
```

As we can see, MSE for AIC model is lowest. So we will choose that. Therefore ARMA(5,2) is the best model and we will forecast with respect to this model.

## Residuals

```
checkresiduals(aic_model)
```

## Residuals from ARIMA(5,0,2) with non-zero mean



```
##
##  Ljung-Box test
##
## data:  Residuals from ARIMA(5,0,2) with non-zero mean
## Q* = 60.265, df = 3, p-value = 5.159e-13
##
## Model df: 8.   Total lags used: 11
```


We can see that residual follows almost a normal distribution which is a good sign for this model.

## Forecasting

```
plot(forecast(aic_model,h=20),type='l')
lines(c((train_size+1):n),test_data, col = 'black')
```



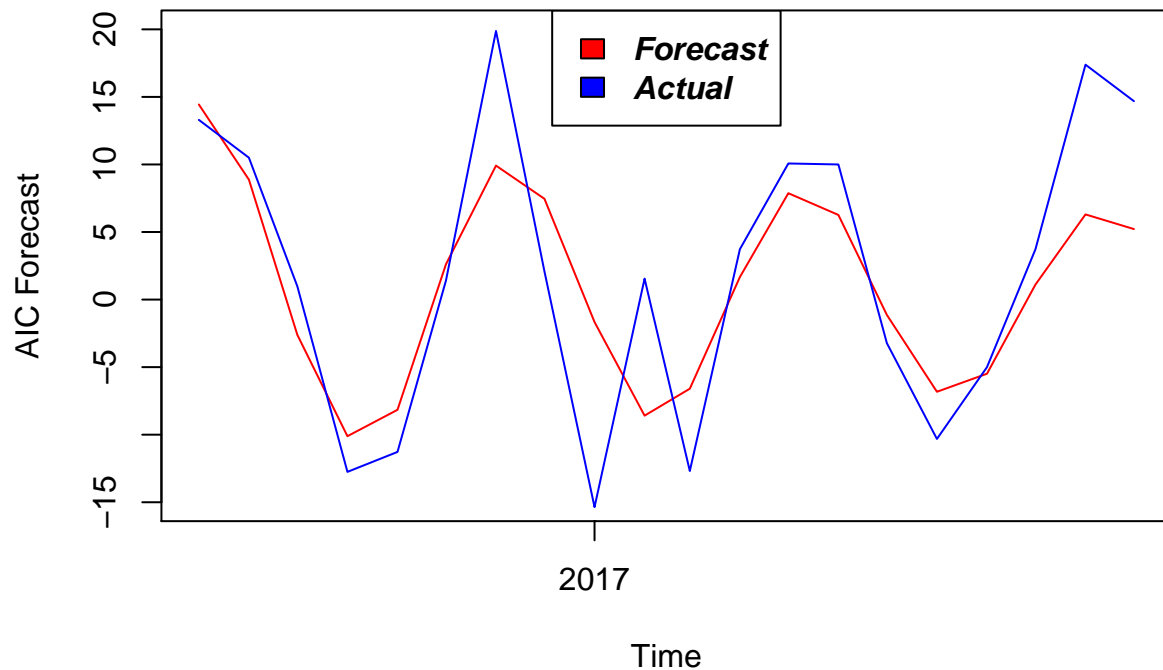
## Forecasts from ARIMA(5,0,2) with non-zero mean



```
plot(Time[(train_size + 1): n],aic_forecast,type='l',col = 'red',
      ylim = c(-15,20),xlab = 'Time',ylab = 'AIC Forecast',
      main = 'Actual Data vs Forecast Data')

lines(Time[(train_size+1):n],test_data, col = 'blue')
legend('top', legend=c("Forecast", "Actual"),fill=c("red", "blue"), text.font=4)
```

## Actual Data vs Forecast Data



## Inverse Transformation

```
y_hat = rep(0, (length(test_data) + 1))
y_hat[1] = data[(length(train_data))]
for (i in 2:(length(test_data) + 1)) {
  y_hat[i] = aic_forecast[i - 1] + y_hat[i - 1]
}

df_compare = data.frame(cbind(y_hat[2:(length(test_data) + 1)], data[c((train_size+2) : length(Time))]))
## Took (train_size + 2), since in the main data we have one extra row, first one.
colnames(df_compare) = c("Predicted", "Actual")
df_compare
```

```
##   Predicted   Actual
## 1  103.76434 104.0375
## 2  112.64052 114.5397
## 3  110.02710 115.5159
## 4   99.91560 102.7637
## 5   91.75646  91.4867
## 6   94.36664  92.8900
## 7  104.28110 112.7694
## 8  111.73824 114.8505
## 9  110.07556  99.4901
## 10 101.48482 101.0396
## 11  94.88823  88.3530
```

```
## 12 96.55488 92.0805
## 13 104.42672 102.1532
## 14 110.69114 112.1538
## 15 109.55664 108.9312
## 16 102.73284 98.6154
## 17 97.24516 93.6137
## 18 98.35416 97.3359
## 19 104.65782 114.7212
## 20 109.87463 129.4048
```

Now let's plot it.

```
plot(Time[(train_size+2):length(Time)],df_compare$Predicted,type='l',col = 'red', ylim = c(80,130),
      ylab = 'Eletric Production',xlab = 'Time',
      main = 'Actual Data vs Forecast Data for Main Data')

lines(Time[(train_size+2):length(Time)],df_compare$Actual, col = 'blue')
legend('bottomright', legend=c("Forecast", "Actual"),fill = c("red", "blue"), text.font=4)
```

